

Lab Assignment no. 2

PRACTICAL NO. 4

NAME: ATHARVA SHANKAR LOKHANDE

PRN:202501070166

BRANCH : ENTC

DIVISION :ET22

SUBJECT : EDS LAB

3.1.1.PANDAS – SERIES CREATION AND MANIPULATION

Write a Python program that takes a list of numbers from the user, creates a Pandas series from it, and then calculates the mean of even and odd numbers separately using the groupby and mean() operations.

Hint: You can group the Series based on whether each number is even or odd.

Input Format:

- The user should enter a list of numbers separated by space when prompted.

Output Format:

- The program should display the mean of even and odd numbers separately.
- Each mean value should be displayed with a label indicating whether it corresponds to even or odd numbers.

Code:

```
import pandas as pd
```

```
# Take inputs from the user to create a list of numbers numbers =
```

```
list(map(int, input().split()))
```

```
# Create a Pandas series from the list of numbers series
```

```
= pd.Series(numbers)
```


```
# Grouping by even and odd numbers and calculating the mean grouped =
```

```
series.groupby(series % 2 == 0).mean()
```

```
# Display the mean of even and odd numbers with labels grouped.index
```

```
= ['Even' if is_even else 'Odd' for is_even in grouped.index] print("Mean
```

```
of even and odd numbers:") print(grouped)
```



4.1.2.DICTIONARY TO DATAFRAME

A dictionary of lists has been provided to you in the editor. Create a DataFrame from the dictionary of lists and perform the listed operations, then display the DataFrame before and after each manipulation.

Create the DataFrame:

- Convert the dictionary to a Pandas DataFrame.

Add a new row:

- Take inputs from the user for the new row data (name, age).
- Add the new row to the DataFrame.
- Display the DataFrame after adding the new row.

Modify a row:

- Modify a specific row by changing the age. Take the row index and new age value from the user.
- Display the DataFrame after modifying the row.

Delete a row:

- **Take the row index to be deleted from the user.**
- **Remove the specified row.**
- **Display the DataFrame after deleting the row.**

Add a new column:

- **Add a column Gender with values taken from the user.**
- **Display the DataFrame after adding the new column.**

Modify a column:

- **Convert names to uppercase.**
- **Display the DataFrame after modifying the column.**

Delete a column:

- **Remove the Age column.**
- **Display the DataFrame after deleting the column.**

Code:

```
import pandas as pd
```

```
# Provided dictionary of lists data
```

```
= {
```

```
    'Name': ['Alice', 'Bob', 'Charlie'],
```

```
    'Age': [25, 30, 35], }
```

```
# Convert the dictionary to a DataFrame df
```

```
= pd.DataFrame(data)
```

```
# Display the original DataFrame print("Original  
DataFrame:") print(df)
```

```
# Adding a new row new_name = input("New  
name: ") new_age = int(input("New age: "))  
df.loc[len(df)]  
= [new_name,new_age]
```

```
# Display the DataFrame after adding a new row print("After adding  
a row:\n",df)
```

```
# Modifying a row row_to_modify = int(input("Index of  
row to modify: ")) new_age_value = int(input("New  
age: ")) df.at[row_to_modify,"Age"] = new_age_value  
# Display the DataFrame after modifying a row  
print("After modifying a row:") print(df)
```

```
# Deleting a row row_to_delete = int(input("Index of row  
to delete: ")) df =  
df.drop(index=row_to_delete).reset_index(drop=True)  
# Display the DataFrame after deleting a row print("After  
deleting a row:") print(df)
```

```
# Adding a new column genders = input("Enter genders separated  
by space: ").split() df["Gender"] = genders
```

```
# Display the DataFrame after adding a new column print("After  
adding a new column:") print(df)
```

```
# Modifying a column df["Name"] =  
df["Name"].str.upper() # Display the DataFrame after  
modifying a column print("After modifying a  
column:") print(df)
```

```
# Deleting a column df = df.drop(columns=["Age"])  
# Display the DataFrame after deleting a column print("After deleting  
a column:")  
print(df)
```

[Home](#)
[Learn Anywhere](#)

202501070137@mitaoe.ac.in

datafram...

```

27 new_age_value = int(input("new age: "))
28 df.at[row_to_modify, "Age"] = new_age_value
# Display the DataFrame after modifying a row

```

Average time

0.071 s

71.50 ms

Maximum time

0.085 s

85.00 ms

1 out of 1 shown test case(s) passed

1 out of 1 hidden test case(s) passed

Test case 1 85ms Debug

Expected output:

Original DataFrame:
..... Name Age
0 ... Alice ... 25
1 ... Bob ... 30
2 ... Charlie ... 35
New name: Susan
New age: 29
After adding a row:
..... Name Age
0 ... Alice ... 25
1 ... Bob ... 30
2 ... Charlie ... 35
3 ... Susan ... 29
Index of row to modify: 0
New age: 27
After modifying a row:
..... Name Age
0 ... Alice ... 27
1 ... Bob ... 30
2 ... Charlie ... 35
3 ... Susan ... 29
Index of row to delete: 3
After deleting a row:
..... Name Age
0 ... Alice ... 27
1 ... Bob ... 30
2 ... Charlie ... 35
Enter genders separated by space: F M M
After adding a new column:
..... Name Age Gender
0 ... Alice ... 27 ... F
1 ... Bob ... 30 ... M
2 ... Charlie ... 35 ... M
After modifying a column:
..... Name Age Gender
0 ... ALICE ... 27 ... F
1 ... BOB ... 30 ... M
2 ... CHARLIE ... 35 ... M
After deleting a column:
..... Name Gender
0 ... ALICE ... F
1 ... BOB ... M
2 ... CHARLIE ... M

Actual output
Original DataFrame:
..... Name Age
0 ... Alice ... 25
1 ... Bob ... 30
2 ... Charlie ... 35
New name: Susan
New age: 29
After adding a row:
..... Name Age
0 ... Alice ... 25
1 ... Bob ... 30
2 ... Charlie ... 35
3 ... Susan ... 29
Index of row to modify: 0
New age: 27
After modifying a row:
..... Name Age
0 ... Alice ... 27
1 ... Bob ... 30
2 ... Charlie ... 35
3 ... Susan ... 29
Index of row to delete: 3
After deleting a row:
..... Name Age
0 ... Alice ... 27
1 ... Bob ... 30
2 ... Charlie ... 35
Enter genders separated by space: F M M
After adding a new column:
..... Name Age Gender
0 ... Alice ... 27 ... F
1 ... Bob ... 30 ... M
2 ... Charlie ... 35 ... M
After modifying a column:
..... Name Age Gender
0 ... ALICE ... 27 ... F
1 ... BOB ... 30 ... M
2 ... CHARLIE ... 35 ... M
After deleting a column:
..... Name Gender
0 ... ALICE ... F
1 ... BOB ... M
2 ... CHARLIE ... M

4.1.3.STUDENT INFORMATION

Write a program to read a text file containing student information (name, age, and grade) using Pandas. Perform the following tasks:

- Display the first five rows of the data frame.
- Calculate the average age of the students (limit the average age up to 2 decimal places).

- **Filter out the students who have a grade above a certain threshold (consider the threshold grade is 'B').**

Note:

Refer to the displayed test cases for better understanding.

Code:

```
import pandas as pd
```

```
# Read the text file into a DataFrame file = input() data = pd.read_csv(file, sep="\s+",  
header=None, names=["Name", "Age", "Grade"])
```

```
print("First five rows:") print(data.head())
```

```
average_age=round(data['Age'].mean(),2)
```

```
print(f"Average age: {average_age}") threshold_grade='B'
```

```
filtered_data=data[data['Grade']<=threshold_grade]
```

```
print(f"Students with a grade up to {threshold_grade}")
```

```
print(filtered_data)
```

studentin...

Submit

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

```

import pandas as pd

# Read the text file into a DataFrame
file = input()
data = pd.read_csv(file, sep="\s+", header=None, names=["Name", "Age", "Grade"])

print("First five rows:")
print(data.head())
average_age=round(data['Age'].mean(),2)
print(f"Average age: {average_age}")
threshold_grade='B'
filtered_data=data[data['Grade']<=threshold_grade]
print(f"Students with a grade up to {threshold_grade}")
print(filtered_data)

```

Average time

0.049 s

49.00 ms

Maximum time

0.075 s

75.00 ms

1 out of 1 shown test case(s) passed

1 out of 1 hidden test case(s) passed

Test case 1

75 ms

Debug

Expected output

Actual output

studentdata.txt

First five rows:

...

0 John 25 A

1 Alin 22 B

2 Emma 24 A

3 Mike 23 C

4 Sony 21 B

Average age: 23.29

Students with a grade up to B:

...

0 John 25 A

1 Alin 22 B

2 Emma 24 A

4 Sony 21 B

5 Amia 26 A

studentdata.txt

First five rows:

...

0 John 25 A

1 Alin 22 B

2 Emma 24 A

3 Mike 23 C

4 Sony 21 B

Average age: 23.29

Students with a grade up to B:

...

0 John 25 A

1 Alin 22 B

2 Emma 24 A

4 Sony 21 B

5 Amia 26 A

Terminal

Test cases

4.2.1.MONTH WITH THE HIGHEST TOTAL SALES

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Group the data by Month and calculate the total sales for each month.

- Find the month with the highest total sales and display it.
- Also, display the total sales for the best month.

Sample Data:

Date	Product	Quantity	Price	City
2025-01-01	Product A	5	20	New York
2025-01-01	Product B	3	15	Los Angeles
2025-01-02	Product A	7	20	New York
2025-01-02	Product C	4	30	Chicago
2025-01-03	Product B	2	15	Chicago
2025-01-03	Product A	8	20	Los Angeles
2025-01-04	Product C	6	30	New York
2025-01-04	Product B	5	15	Los Angeles
2025-01-05	Product A	3	20	Chicago
2025-01-05	Product C	10	30	Los Angeles

Note:

The data cannot be displayed in the file. You can refer to the sample data provided for insights.

Code:

```
import pandas as pd
```

```
# Prompt the user for the file name file_name
```

```
= input()
```

```
# Load the data df = pd.read_csv(file_name)
df['Total_Sales']=df['Quantity']*df['Price']
df['Date']=pd.to_datetime(df['Date'])
df['Month']=df['Date'].dt.to_period('M')
monthly_sales=df.groupby('Month')['Total_
Sales'].sum() # Find the month with the
highest total sales best_month =
monthly_sales.idxmax() highest_sales =
monthly_sales.max()

print(f"Best month: {best_month}") print(f"Total sales:
${highest_sales:.2f}")
```

```

1 import pandas as pd
2
3 # Prompt the user for the file name
4 file_name = input()
5
6 # Load the data
7 df = pd.read_csv(file_name)
8
9
10 df['Total_Sales'] = df['Quantity'] * df['Price']
11 df['Date'] = pd.to_datetime(df['Date'])
12 df['Month'] = df['Date'].dt.to_period('M')
13 monthly_sales = df.groupby('Month')['Total_Sales'].sum()
14 # Find the month with the highest total sales
15 best_month = monthly_sales.idxmax()
16 highest_sales = monthly_sales.max()
17
18 print(f"Best month: {best_month}")
19 print(f"Total sales: ${highest_sales:.2f}")
20

```

Average time: 0.033 s
 Maximum time: 0.067 s
 1 out of 1 shown test case(s) passed
 2 out of 2 hidden test case(s) passed

Test case 1 (67 ms) [Debug] [Test]

Expected output	Actual output
sales_data.csv	sales_data.csv
Best month: 2025-01	Best month: 2025-01
Total sales: \$1210.00	Total sales: \$1210.00

Terminal | Test cases

4.2.2.BEST SELLING PRODUCT

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Find the product that sold the most in terms of quantity sold.
- Display the product that sold the most and the total quantity sold for that product.

Sample Data:

Date	Product	Quantity	Price	City
2025-01-01	Product A	5	20	New York

2025-01-01	Product B	3	15	Los Angeles
2025-01-02	Product A	7	20	New York
2025-01-02	Product C	4	30	Chicago
2025-01-03	Product B	2	15	Chicago
2025-01-03	Product A	8	20	Los Angeles
2025-01-04	Product C	6	30	New York
2025-01-04	Product B	5	15	Los Angeles
2025-01-05	Product A	3	20	Chicago
2025-01-05	Product C	10	30	Los Angeles

Note:

The data cannot be displayed in the file. You can refer to the sample data provided for insights.

Code:

```
import pandas as pd
```

```
# Prompt the user for the file name file_name
```

```
= input()
```

```
# Load the data df = pd.read_csv(file_name)
```

```
df['Total_Sales']=df['Quantity']*df['Price']
```

```
product_sales=df.groupby('Product')['Quan
```

```
tity'].sum()
```

```
# Find the product with the highest total quantity sold best_product
```

```
=product_sales.idxmax() highest_quantity = product_sales.max()
```

```
# Display the result print(f"Best selling  
product:      {best_product}") print(f"Total  
            quantity      sold:  
{highest_quantity}")
```

```

1 import pandas as pd
2
3 # Prompt the user for the file name
4 file_name = input()
5
6 # Load the data
7 df = pd.read_csv(file_name)
8
9
10 df['Total_Sales'] = df['Quantity'] * df['Price']
11 product_sales = df.groupby('Product')['Total_Sales'].sum()
12
13 # Find the product with the highest total quantity sold
14 best_product = product_sales.idxmax()
15 highest_quantity = product_sales.max()
16
17 # Display the result
18 print(f"Best selling product: {best_product}")
19 print(f"Total quantity sold: {highest_quantity}")
20

```

Average time: 0.026 s
 Maximum time: 0.055 s
 1 out of 1 shown test case(s) passed
 2 out of 2 hidden test case(s) passed

Test case 1 55 ms Debug

Expected output	Actual output
sales_data.csv	sales_data.csv
Best selling product: Product A	Best selling product: Product A
Total quantity sold: 23	Total quantity sold: 23

Terminal Test cases

4.2.3.CITY THAT SOLD THE MOST PRODUCT

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Group the data by City and calculate the total quantity of products sold for each city.
- Find the city that sold the most products (based on the total quantity sold).

Sample Data:

<u>Date</u>	<u>Product</u>	<u>Quantity</u>	<u>Price</u>	<u>City</u>
<u>2025-01-01</u>	<u>Product A</u>	<u>5</u>	<u>20</u>	<u>New York</u>
<u>2025-01-01</u>	<u>Product B</u>	<u>3</u>	<u>15</u>	<u>Los Angeles</u>

<u>2025-01-02</u>	<u>Product A</u>	<u>7</u>	<u>20</u>	<u>New York</u>
<u>2025-01-02</u>	<u>Product C</u>	<u>4</u>	<u>30</u>	<u>Chicago</u>
<u>2025-01-03</u>	<u>Product B</u>	<u>2</u>	<u>15</u>	<u>Chicago</u>
<u>2025-01-03</u>	<u>Product A</u>	<u>8</u>	<u>20</u>	<u>Los Angeles</u>
<u>2025-01-04</u>	<u>Product C</u>	<u>6</u>	<u>30</u>	<u>New York</u>
<u>2025-01-04</u>	<u>Product B</u>	<u>5</u>	<u>15</u>	<u>Los Angeles</u>
<u>2025-01-05</u>	<u>Product A</u>	<u>3</u>	<u>20</u>	<u>Chicago</u>
<u>2025-01-05</u>	<u>Product C</u>	<u>10</u>	<u>30</u>	<u>Los Angeles</u>

Note:

The data cannot be displayed in the file. You can refer to the sample data provided for insights.

Code:

```
import pandas as pd
```

```
# Prompt the user for the file name file_name
```

```
= input()
```

```
# Load the data df = pd.read_csv(file_name)
```

```
city_totals=df.groupby('City
```

```
')['Quantity'].sum()
```

```
best_city=city_totals.idxma
```

```
x()
```

```
# write the code..
```

```
# Display the result print(f"City sold the most products:
{best_city}")
```

The screenshot shows a code editor with a Python program that reads a CSV file, groups data by city, and prints the city with the most products. Below the code editor, the test results are displayed, showing that 1 out of 1 shown test case(s) passed and 2 out of 2 hidden test case(s) passed. The test case details show the expected output as 'sales_data.csv' and 'City sold the most products: Los Angeles', which matches the actual output.

```
1 import pandas as pd
2
3 # Prompt the user for the file name
4 file_name = input()
5
6 # Load the data
7 df = pd.read_csv(file_name)
8
9 city_totals=df.groupby('City')['Quantity'].sum()
10 best_city=city_totals.idxmax()
11 # write the code..
12
13 # Display the result
14 print(f"City sold the most products: {best_city}")
15
```

Average time: 0.024 s, Maximum time: 0.049 s
 1 out of 1 shown test case(s) passed
 2 out of 2 hidden test case(s) passed

Test case 1: 49 ms, Debug, Test cases

Expected output: sales_data.csv, City sold the most products: Los Angeles
 Actual output: sales_data.csv, City sold the most products: Los Angeles

4.2.4.MOST FREQUENTLY SOLD PRODUCT PAIRS

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the following columns: Date, Product, Quantity, Price, and City.
- For each date, find all pairs of products that were sold together (i.e., two products sold on the same date).
- Output the product pair/s that was sold most frequently.

Sample Data:

Date	Product	Quantity	Price	City
------	---------	----------	-------	------

2025-01-01	Product A	5	20	New York
2025-01-01	Product B	3	15	Los Angeles
2025-01-02	Product A	7	20	New York
2025-01-02	Product C	4	30	Chicago
2025-01-03	Product B	2	15	Chicago
2025-01-03	Product A	8	20	Los Angeles
2025-01-04	Product C	6	30	New York
2025-01-04	Product B	5	15	Los Angeles
2025-01-05	Product A	3	20	Chicago
2025-01-05	Product C	10	30	Los Angeles

Explanation:

Transactions:

- 2025-01-01: Product A, Product B
- 2025-01-02: Product A, Product C
- 2025-01-03: Product B, Product A • 2025-01-04: Product C, Product B
- 2025-01-05: Product A, Product C

Now, let's count how often the pairs of products appear together:

- Product A and Product B: Appear in transactions on 2025-01-01 and 2025-01-03.
- Product A and Product C: Appear in transactions on 2025-01-02 and 2025-01-05.
- Product B and Product C: Appears in transactions on 2025-01-04.

Most Frequent Product Combinations:

- Product A and Product B (2 times)
- Product A and Product C (2 times)

Note:

The data cannot be displayed in the file. You can refer to the sample data provided for insights.

```
Code: import pandas as pd from
itertools import combinations from
collections import Counter
```

```
# Prompt user to input the file name file_name
= input()
```

```
# Read data from the specified CSV file df
= pd.read_csv(file_name)
```

```
product_pairs=[] for date,group in df.groupby('Date'):
    products=group['Product'].unique()
    if len(products)>1:
        pairs=combinations(sorted(products),2)
    product_pairs.extend(pairs)
```

```
pair_counter=Counter(product_pairs)
```

```
max_freq=max(pair_counter.values()) most_frequent_pairs=[pair for pair ,count in
pair_counter.items() if count==max_freq]
```

```
for pair in most_frequent_pairs:    print(f"{pair[0]} and
{pair[1]}: {max_freq} times")
```

```
1 import pandas as pd
2 from itertools import combinations
3 from collections import Counter
4
5 # Prompt user to input the file name
6 file_name = input()
7
8 # Read data from the specified CSV file
9 df = pd.read_csv(file_name)
10
11 product_pairs=[]
12 for date,group in df.groupby('Date'):
13     products=group['Product'].unique()
14     if len(products)>1:
15         pairs=combinations(sorted(products),2)
16         product_pairs.extend(pairs)
17
18 pair_counter=Counter(product_pairs)
19
20 max_freq=max(pair_counter.values())
21 most_frequent_pairs=[pair for pair ,count in pair_counter.items() if count==max_freq]
22
23 for pair in most_frequent_pairs:
24     print(f"{pair[0]} and {pair[1]}: {max_freq} times")
25
26 # Output the most frequent product pairs
```

Average time: 0.026 s (26.00 ms) | Maximum time: 0.050 s (50.00 ms)

1 out of 1 shown test case(s) passed
2 out of 2 hidden test case(s) passed

Test case 1 (50 ms) [Debug]

Expected output	Actual output
sales_data.csv	sales_data.csv
Product A and Product B: 2 times	Product A and Product B: 2 times
Product A and Product C: 2 times	Product A and Product C: 2 times

Terminal | Test cases

4.2.5.TITANIC DATASET ANALYSIS AND DATA CLEANING

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset. For each question, perform necessary data cleaning, transformations, and calculations as required.

1. Display the first 5 rows of the dataset.

2. Display the last 5 rows of the dataset.
3. Get the shape of the dataset (number of rows and columns).
4. Get a summary of the dataset (using `.info()`).
5. Get basic statistics (mean, standard deviation, etc.) of the dataset using `.describe()`.
6. Check for missing values and display the count of missing values for each column.
7. Fill missing values in the 'Age' column with the median age.
8. Fill missing values in the 'Embarked' column with the most frequent value (`mode()`).
9. Drop the 'Cabin' column due to many missing values.
10. Create a new column, 'FamilySize' by adding the 'SibSp' and 'Parch' columns.

The Titanic dataset contains columns as shown below,

Sample Data:

PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked

1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S

2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC
17599,71.2833,C85,C

PassengerId	Survived	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
-------------	----------	------	-----	-----	-------	-------	--------	------	-------	----------

3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S
6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S
10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C

Code: import pandas as pd

import numpy as np

Load the Titanic dataset data = pd.read_csv('Titanic-Dataset.csv')

1. Display the first 5 rows of the dataset print(data.head(5))

2. Display the last 5 rows of the dataset print(data.tail(5))

3. Get the shape of the dataset print(data.shape)

4. Get a summary of the dataset (info) print(data.info())

5. Get basic statistics of the dataset print(data.describe())

6. Check for missing values

```
print(data.isnull().sum())
```

Explorer	TitanicDat...	0	Submit
Average time	0.200 s	Maximum time	0.200 s
200.00 ms		200.00 ms	
1 out of 1 shown test case(s) passed			
Expected output	Actual output		
PassengerId Survived Pclass ... Fare Cabin Embarked	PassengerId Survived Pclass ... Fare Cabin Embarked		
0 ... 1 ... 0 ... 3 ... 7.2500 NaN S	0 ... 1 ... 0 ... 3 ... 7.2500 NaN S		
1 ... 2 ... 1 ... 1 ... 71.2833 C85 C	1 ... 2 ... 1 ... 1 ... 71.2833 C85 C		
2 ... 3 ... 1 ... 3 ... 7.9250 NaN S	2 ... 3 ... 1 ... 3 ... 7.9250 NaN S		
3 ... 4 ... 1 ... 1 ... 53.1000 C123 S	3 ... 4 ... 1 ... 1 ... 53.1000 C123 S		
4 ... 5 ... 0 ... 3 ... 8.0500 NaN S	4 ... 5 ... 0 ... 3 ... 8.0500 NaN S		
[5 rows x 12 columns]	[5 rows x 12 columns]		
PassengerId Survived Pclass ... Fare Cabin Embarked	PassengerId Survived Pclass ... Fare Cabin Embarked		
886 ... 887 ... 0 ... 2 ... 13.00 NaN S	886 ... 887 ... 0 ... 2 ... 13.00 NaN S		
887 ... 888 ... 1 ... 1 ... 30.00 842 S	887 ... 888 ... 1 ... 1 ... 30.00 842 S		
888 ... 889 ... 0 ... 3 ... 23.45 NaN S	888 ... 889 ... 0 ... 3 ... 23.45 NaN S		
889 ... 890 ... 1 ... 1 ... 30.00 C148 C	889 ... 890 ... 1 ... 1 ... 30.00 C148 C		
890 ... 891 ... 0 ... 3 ... 7.75 NaN Q	890 ... 891 ... 0 ... 3 ... 7.75 NaN Q		
[5 rows x 12 columns]	[5 rows x 12 columns]		
(891, 12)	(891, 12)		
<class 'pandas.core.frame.DataFrame'>	<class 'pandas.core.frame.DataFrame'>		
RangeIndex: 891 entries, 0 to 890	RangeIndex: 891 entries, 0 to 890		
Data columns (total 11 columns):	Data columns (total 12 columns):		
# Column Non-Null Count Dtype	# Column Non-Null Count Dtype		
0 PassengerId 891 non-null int64	0 PassengerId 891 non-null int64		
1 Survived 891 non-null int64	1 Survived 891 non-null int64		
2 Pclass 891 non-null int64	2 Pclass 891 non-null int64		
3 Name 891 non-null object	3 Name 891 non-null object		
4 Sex 891 non-null object	4 Sex 891 non-null object		
5 Age 714 non-null float64	5 Age 714 non-null float64		
6 SibSp 891 non-null int64	6 SibSp 891 non-null int64		
7 Parch 891 non-null int64	7 Parch 891 non-null int64		
8 Ticket 891 non-null object	8 Ticket 891 non-null object		
9 Fare 891 non-null float64	9 Fare 891 non-null float64		
10 Cabin 204 non-null object	10 Cabin 204 non-null object		
11 Embarked 889 non-null object	11 Embarked 889 non-null object		
dtypes: float64(2), int64(5), object(5)	dtypes: float64(2), int64(5), object(5)		
memory usage: 66.2+ KB	memory usage: 66.2+ KB		
None	None		
Terminal	Test cases		

PassengerId Survived Pclass ... SibSp Parch	PassengerId Survived Pclass ... SibSp Parch
Fare	Fare
count 891.000000 891.000000 891.000000 ... 891.000000 891.000000 891.000000 89	count 891.000000 891.000000 891.000000 ... 891.000000 891.000000 891.000000 89
1.000000	1.000000
mean 445.000000 0.383838 2.308542 ... 0.523000 0.381594 3	mean 445.000000 0.383838 2.308542 ... 0.523000 0.381594 3
2.284288	2.284288
std 257.353842 0.485592 0.836071 ... 1.182743 0.986057 4	std 257.353842 0.485592 0.836071 ... 1.182743 0.986057 4
9.693429	9.693429
min 1.000000 0.000000 1.000000 ... 0.000000 0.000000 0	min 1.000000 0.000000 1.000000 ... 0.000000 0.000000 0
0.000000	0.000000
25% 223.500000 0.000000 2.000000 ... 0.000000 0.000000 0	25% 223.500000 0.000000 2.000000 ... 0.000000 0.000000 0
7.918400	7.918400
50% 445.000000 0.000000 3.000000 ... 0.000000 0.000000 1	50% 445.000000 0.000000 3.000000 ... 0.000000 0.000000 1
4.454300	4.454300
75% 668.500000 1.000000 3.000000 ... 1.000000 0.000000 3	75% 668.500000 1.000000 3.000000 ... 1.000000 0.000000 3
1.000000	1.000000
max 891.000000 1.000000 3.000000 ... 8.000000 6.000000 51	max 891.000000 1.000000 3.000000 ... 8.000000 6.000000 51
2.329100	2.329100
[8 rows x 7 columns]	[8 rows x 7 columns]
PassengerId 0	PassengerId 0
Survived 0	Survived 0
Pclass 0	Pclass 0
Name 0	Name 0
Sex 0	Sex 0
Age 177	Age 177
SibSp 0	SibSp 0
Parch 0	Parch 0
Ticket 0	Ticket 0
Fare 0	Fare 0
Cabin 687	Cabin 687
Embarked 2	Embarked 2
dtype: int64	dtype: int64
Terminal	Test cases

4.2.6.TITANIC DATASET ANALYSIS AND DATA CLEANING – 2

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Create a new column 'IsAlone' which is 1 if the passenger is alone (FamilySize = 0), otherwise 0.
2. Convert the 'Sex' column to numeric values (male: 0, female: 1).

3. One-hot encode the 'Embarked' column, dropping the first category.
4. Get the mean age of passengers.
5. Get the median fare of passengers.
6. Get the number of passengers by class.
7. Get the number of passengers by gender.
8. Get the number of passengers by survival status.
9. Calculate the survival rate of passengers.
10. Calculate the survival rate by gender.

The Titanic dataset contains columns as shown below,

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked

Sample Data:

PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked

1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S

2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC
17599,71.2833,C85,C

3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S

4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S

5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S

6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q

7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S

8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S

9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina
Berg)",female,27,0,2,347742,11.1333,,S

10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C

Code: import pandas as pd

import numpy as np

Load the Titanic dataset data =

pd.read_csv('Titanic-Dataset.csv') data['FamilySize']

= data['SibSp'] + data['Parch']

1. Create a new column 'IsAlone' (1 if alone, 0 otherwise)

data['IsAlone']=np.where(data['FamilySize']>0,0,1)

2. Convert 'Sex' to numeric (male: 0, female: 1)

data['Sex']=data['Sex'].map({'male':0,'female':1})

3. One-hot encode the 'Embarked' column

data=pd.get_dummies(data,columns=['Embarked'],drop_first=True)

4. Get the mean age of passengers print(data['Age'].mean())

5. Get the median fare of passengers print(data['Fare'].median())

6. Get the number of passengers by cla print(data['Pclass'].value_counts()) # 7. Get the number of passengers by gender print(data['Sex'].value_counts())

8. Get the number of passengers by survival status `print(data['Survived'].value_counts())` # 9. Calculate the survival rate `total=len(data) survivals=data['Survived'].sum() print((survivals/total))`

10. Calculate the survival rate by gender `print(data.groupby('Sex')['Survived'].mean())`

```

1 import pandas as pd
2 import numpy as np
3
4 # Load the Titanic dataset
5 data = pd.read_csv('Titanic-Dataset.csv')
6 data['FamilySize'] = data['SibSp'] + data['Parch']
7
8
9 # 1. Create a new column 'IsAlone' (1 if alone, 0 otherwise)
10 data['IsAlone'] = np.where(data['FamilySize'] > 0, 0, 1)
11
12 # 2. Convert 'Sex' to numeric (male: 0, female: 1)
13 data['Sex'] = data['Sex'].map({'male': 0, 'female': 1})
14
15 # 3. One-hot encode the 'Embarked' column
16 data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)
17
18 # 4. Get the mean age of passengers
19 print(data['Age'].mean())

```

Average time: 0.084 s
Maximum time: 0.084 s
84.00 ms 84.00 ms

1 out of 1 shown test case(s) passed

Test case 1 34 ms Debug

Expected output	Actual output
29.69911764705882	29.69911764705882
14.4542	14.4542
3 --- 491	3 --- 491
1 --- 216	1 --- 216
2 --- 184	2 --- 184
Name: Pclass, dtype: int64	Name: Pclass, dtype: int64
0 --- 577	0 --- 577
1 --- 314	1 --- 314
Name: Sex, dtype: int64	Name: Sex, dtype: int64
0 --- 549	0 --- 549
1 --- 342	1 --- 342
Name: Survived, dtype: int64	Name: Survived, dtype: int64
0.3838383838383838	0.3838383838383838
Sex	Sex
0 --- 0.188908	0 --- 0.188908
1 --- 0.742038	1 --- 0.742038
Name: Survived, dtype: float64	Name: Survived, dtype: float64

Terminal Test cases

4.2.7. TITANIC DATASET ANALYSIS AND DATA CLEANING – 3

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Calculate the survival rate by class.
2. Calculate the survival rate by embarkation location (Embarked_S).
3. Calculate the survival rate by family size (FamilySize).
4. Calculate the survival rate by being alone (IsAlone).
5. Get the average fare by passenger class (Pclass).
7. Get the average age by survival status (Survived).
8. Get the average fare by survival status (Survived).
9. Get the number of survivors by class (Pclass).
10. Get the number of non-survivors by class (Pclass).

The Titanic dataset contains columns as shown below,

6. Get the average age by passenger class (Pclass).

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked

Sample Data:

PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked

1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S

2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC
17599,71.2833,C85,C

3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S

4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S

5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S

6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S
10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C

Code: import pandas as pd

import numpy as np

```
# Load the Titanic dataset data = pd.read_csv('Titanic-Dataset.csv') data['FamilySize']  
= data['SibSp'] + data['Parch'] data['IsAlone'] = np.where(data['FamilySize'] > 0, 0, 1)  
data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)
```

```
# 1. Calculate the survival rate by class print(data.groupby('Pclass')['Survived'].mean())
```

```
# 2. Calculate the survival rate by embarked location
```

```
print(data.groupby('Embarked_S')['Survived'].mean()) # 3. Calculate the survival rate by family  
size print(data.groupby('FamilySize')['Survived'].mean())
```

```
# 4. Calculate the survival rate by being alone print(data.groupby('IsAlone')['Survived'].mean())
```

```
# 5. Get the average fare by class print(data.groupby('Pclass')['Fare'].mean())
```

```
# 6. Get the average age by class print(data.groupby('Pclass')['Age'].mean())
```

```
# 7. Get the average age by survival status print(data.groupby('Survived')['Age'].mean())
```


```
# 8. Get the average fare by survival status print(data.groupby('Survived')['Fare'].mean())
```

9. Get the number of survivors by class

```
print(data[data['Survived']==1]['Pclass'].value_counts())
```

10. Get the number of non-survivors by class

```
print(data[data['Survived']==0]['Pclass'].value_counts())
```



4.2.8. TITANIC DATASET ANALYSIS AND DATA CLEANING – 4

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Get the number of survivors by gender (Sex).
2. Get the number of non-survivors by gender (Sex).
3. Get the number of survivors by embarkation location (Embarked_S).
4. Get the number of non-survivors by embarkation location (Embarked_S).
5. Calculate the percentage of children (Age < 18) who survived.
6. Calculate the percentage of adults (Age >= 18) who survived.
7. Get the median age of survivors.
8. Get the median age of non-survivors.
9. Get the median fare of survivors.

10. Get the median fare of non-survivors.

The Titanic dataset contains columns as shown below,

Sample Data:

PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, Ticket, Fare, Cabin, Embarked

1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S

2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC
17599,71.2833,C85,C

3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S

4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S

5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S

6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q

7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S

8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S

9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina
Berg)",female,27,0,2,347742,11.1333,,S

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked

10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C

Code: import pandas as pd

import numpy as np

Load the Titanic dataset data = pd.read_csv('Titanic-Dataset.csv') data

= pd.get_dummies(data, columns=['Embarked'], drop_first=True)

#1. Get the number of survivors by gender

```
print(data[data['Survived']==1] ['Sex'].value_counts())
```

#2. Get the number of non-survivors by gender

```
print(data[data['Survived']==0] ['Sex'].value_counts())
```

#3. Get the number of survivors by embarked location

```
print(data[data['Survived']==1] ['Embarked_S'].value_counts())
```

#4. Get the number of non-survivors by embarked location `print(data[data['Survived']==0] ['Embarked_S'].value_counts())`

```
print(data[data['Age']<18] ['Survived'].mean())
```

#6. Calculate the percentage of adults (Age >= 18) who survived

```
print(data[data['Age']>=18] ['Survived'].mean())
```

```
print(data[data['Survived']==1] ['Age'].median())
```

#8. Get the median age of non-survivors

```
print(data[data['Survived']==0] ['Age'].median())
```

```
print(data[data['Survived']==1] ['Fare'].median())
```

#10. Get the median fare of non-survivors

```
print(data[data['Survived']==0] ['Fare'].median())
```

The screenshot displays a Jupyter Notebook environment with a file explorer on the left showing 'titanicDat...'. The main area contains a Python script with 32 lines of code. The code imports pandas and numpy, loads the Titanic dataset, and performs several data analysis tasks: counting survivors and non-survivors by gender and embarked location, calculating the mean survival rate for adults, and finally, calculating the median fare for non-survivors. The execution bar at the bottom indicates that the code ran successfully, with an average time of 0.086 seconds and a maximum time of 86.00 milliseconds. A test case summary shows '1 out of 1 shown test case(s) passed'.

```
1 import pandas as pd
2 import numpy as np
3
4 # Load the Titanic dataset
5 data = pd.read_csv('Titanic-Dataset.csv')
6 data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)
7
8
9 #1. Get the number of survivors by gender
10 print(data[data['Survived']==1] ['Sex'].value_counts())
11
12 #2. Get the number of non-survivors by gender
13 print(data[data['Survived']==0] ['Sex'].value_counts())
14
15 #3. Get the number of survivors by embarked location
16 print(data[data['Survived']==1] ['Embarked_S'].value_counts())
17
18 #4. Get the number of non-survivors by embarked location
19 print(data[data['Survived']==0] ['Embarked_S'].value_counts())
20
21 print(data[data['Age']<18] ['Survived'].mean())
22
23 #6. Calculate the percentage of adults (Age >= 18) who survived
24 print(data[data['Age']>=18] ['Survived'].mean())
25
26 print(data[data['Survived']==1] ['Age'].median())
27
28
29
30
31
32
```

Average time: 0.086 s
Maximum time: 0.086 s
1 out of 1 shown test case(s) passed

Test case 1	Expected output	Actual output
female	233	female: 233
male	189	male: 189
Name: Sex, dtype: int64		Name: Sex, dtype: int64
female	468	female: 468
male	81	male: 81
Name: Sex, dtype: int64		Name: Sex, dtype: int64
1	217	1: 217
0	125	0: 125
Name: Embarked_S, dtype: int64		Name: Embarked_S, dtype: int64
1	427	1: 427
0	122	0: 122
Name: Embarked_S, dtype: int64		Name: Embarked_S, dtype: int64
0.539823088495575		0.539823088495575
0.3810316139767055		0.3810316139767055
26.0		26.0
26.0		26.0
26.0		26.0
10.5		10.5

